

Domain-specific acceleration and auto-parallelization of legacy scientific code in FORTRAN 77 using source-to-source compilation

Wim Vanderbauwhede*, Gavin Davidson

School of Computing Science, University of Glasgow, Glasgow, Scotland



ARTICLE INFO

Article history:

Received 8 February 2017

Revised 25 May 2018

Accepted 4 June 2018

Available online 12 June 2018

Keywords:

Fortran

GPGPU

OpenCL

Source-to-source compilation

Auto-parallelization

Acceleration

ABSTRACT

Massively parallel accelerators such as GPGPUs, manycores and FPGAs represent a powerful and affordable tool for scientists who look to speed up simulations of complex systems. However, porting code to such devices requires a detailed understanding of heterogeneous programming tools and effective strategies for parallelization. In this paper we present a source to source compilation approach with whole-program analysis to automatically transform single-threaded FORTRAN 77 legacy code into OpenCL-accelerated programs with parallelized kernels.

The main contributions of our work are: (1) whole-source refactoring to allow any subroutine in the code to be offloaded to an accelerator. (2) Minimization of the data transfer between the host and the accelerator by eliminating redundant transfers. (3) Pragmatic auto-parallelization of the code to be offloaded to the accelerator by identification of parallelizable *maps* and *reductions*.

We have validated the code transformation performance of the compiler on the NIST FORTRAN 78 test suite and several real-world codes: the Large Eddy Simulator for Urban Flows, a high-resolution turbulent flow model; the shallow water component of the ocean model Gmodel; the Linear Baroclinic Model, an atmospheric climate model and Flexpart-WRF, a particle dispersion simulator.

The automatic parallelization component has been tested on as 2-D Shallow Water model (2DSW) and on the Large Eddy Simulator for Urban Flows (UFLES) and produces a complete OpenCL-enabled code base. The fully OpenCL-accelerated versions of the 2DSW and the UFLES are resp. 9x and 20x faster on GPU than the original code on CPU, in both cases this is the same performance as manually ported code.

© 2018 The Authors. Published by Elsevier Ltd.

This is an open access article under the CC BY license. (<http://creativecommons.org/licenses/by/4.0/>)

1. Background

A large amount of scientific code (both “legacy” code and new code) is still effectively written in FORTRAN 77. Fig. 1 shows the relative citations (citations per revision normalized to sum of citations for all revisions) for Google Scholar and ScienceDirect for each of the main revisions of Fortran. We collected results for the past 10 years (2006–2016) and also since the release of FORTRAN 77 (1978–2016). As an absolute reference, there were 15,700 citations in Google Scholar mentioning FORTRAN 77 between 2006 and 2016. It is clear that FORTRAN 77 is still widely used and that the latest standards (2003, 2008) have not yet found widespread adoption.

Based on the above evidence – and also on our own experience of collaboration with scientists – the current state of affairs

is that for many scientists, FORTRAN 77 is still the language of choice for writing models. There is also a vast amount of legacy code in FORTRAN 77. Because the FORTRAN 77 language was designed with assumptions and requirements very different from today's, code written in it has inherent issues with readability, scalability, maintainability and parallelization. A comprehensive discussion of the issues can be found in [1]. As a result, many efforts have been aimed at refactoring legacy code, either interactive or automatic, and to address one or several of these issues.

Our work is part of that effort, but we are specifically interested in automatically *refactoring Fortran for OpenCL-based accelerators*. In this paper we present a source compilation approach to transform sequential FORTRAN 77 legacy code into high-performance OpenCL-accelerated programs with auto-parallelized kernels without need for directives or extra information from the user.

2. Heterogeneous computing and accelerators

By heterogeneous computing we mean computing on a system comprising a (multicore) host processor and an accelerator,

* Corresponding author.

E-mail addresses: wim.vanderbauwhede@glasgow.ac.uk, wim@dcs.gla.ac.uk (W. Vanderbauwhede).

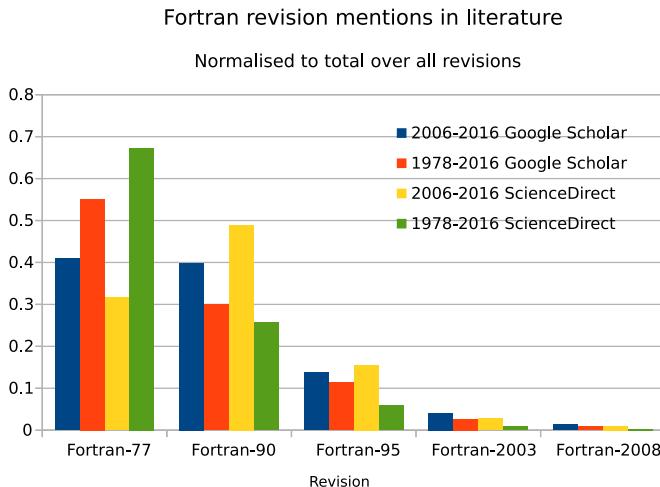


Fig. 1. Literature mentions of different revisions of Fortran using Google Scholar and ScienceDirect.

e.g. a GPGPU, FPGA or a manycore device such as the Intel Xeon Phi. Many scientific codes have already been investigated for and ported manually to GPUs, and excellent performance benefits have been reported. There are many approaches to programming accelerators, but we restrict our discussion to open standards and do not discuss commercial solutions tied to a particular vendor or platform; and we will only discuss solutions that work in Fortran.

2.1. OpenCL

The OpenCL framework [2] presents an abstraction of the accelerator hardware based on the concept of *host* and *device*. A programmer writes one or more *kernels* that are run directly by the accelerator and a *host program* that is run on the system's main CPU. The host program handles memory transfers to the device and initializing computations and the kernels do the bulk of the processing, in parallel on the device. The main advantage of OpenCL over proprietary solutions such as e.g. CUDA (to which it is very similar) is that it supported by a wide range of devices, including multicore CPUs, FPGAs and GPUs. From the programmer perspective, OpenCL is very flexible but quite low level and requires a lot of boilerplate code to be written. This is a considerable barrier for adoption by scientists. Furthermore, there is no official Fortran support for OpenCL: the host API is C/C++, the kernel language is based on a subset of C99. To remedy this we have developed [3] a Fortran API for OpenCL.¹

2.2. OpenACC and OpenMP

OpenACC² takes a directive based approach to heterogeneous programming that affords a higher level of abstraction for parallel programming than OpenCL or CUDA. In a basic example, a programmer adds *pragmas* (compiler directives) to the original (sequential) code to indicate which parts of the code are to be accelerated. The new source code, including directives, is then processed by the OpenACC compiler and programs that can run on accelerators are produced. There are a number of extra directives that allow for optimization and tuning to allow for the best possible performance.

With OpenMP version 4, the popular OpenMP standard³ for shared-memory parallel programming now also supports acceler-

ators. The focus of both standards is slightly different, the main difference being that OpenMP allows conventional OpenMP directives to be combined with accelerator directives, whereas OpenACC directives are specifically designed for offloading computation to accelerators.

Both these annotation-based approaches are local: they deal with parallelization of relatively small blocks and are not aware of the whole code base, and this makes them both harder to use and less efficient. To use either on legacy FORTRAN 77 code, it is not enough to insert the pragmas: the programmer has to ensure that the code to be offloaded is free of global variables, which means complete removal of all common block variables or providing a list of shared variables as annotation. The programmer must also think carefully about the data movement between the host and the device, otherwise performance is poor.

2.3. Raising the abstraction level

Our approach allows an even higher level of abstraction than that offered by OpenACC or OpenMP: the programmer does not need to consider how to achieve program parallelization, but only to mark (using a single annotation) which subroutines will be parallelized and offloaded to the accelerator. Our compiler provides a fully automatic conversion of a complete FORTRAN 77 codebase to Fortran 95 with OpenCL kernels. Consequently, the scientists can keep writing the code in FORTRAN 77, and the original code base is always intact.

3. Existing source-to-source compilers and refactoring tools

A conventional compiler consumes source code and produces binaries. A source-to-source compiler produces transformed source code from the original source. This transformation can be e.g. refactoring, parallelization or translation to a different language. The advantage is that the resulting code can be modified by the programmer if desired and compiled with a compiler of choice.

There are a number of source-to-source compilers and refactoring tools for Fortran available. However, very few of them actually support FORTRAN 77. The most well known are the ROSE framework⁴ from LLNL [4], which relies on the Open Fortran Parser (OFP).⁵ This parser claims to support the Fortran 2008 standard. Furthermore, there is the language-fortran⁶ parser which claims to support FORTRAN 77 to Fortran 2003. A refactoring framework which claims to support FORTRAN 77 is CamFort [5], according to its documentation it supports Fortran 66, 77, and 90 with various legacy extensions.

We tested OFP 0.8.3, language-fortran 0.5.1 and CamFort 0.804 using the NIST FORTRAN 78 test suite (discussed in more detail in Section 5). All three parsers failed to parse any of the provided sources. Consequently we could not use either of these as a starting point.

Like CamFort, the Eclipse-based interactive refactoring tool Photon [6], which supports FORTRAN 77 - 2008, is not a whole-source compiler, but works on a per-file basis (which is in fact what most compilers do). Both CamFort and Photon provide very useful refactorings, but these are limited to the scope of a code unit. For effective refactoring of common blocks, and determination of data movement direction, as well as for effective acceleration, whole-source code (inter-procedural) analysis and refactoring is essential.

A long-running project which does support inter-procedural analysis is PIPS⁷, started in the 1990's. The PIPS tool does sup-

¹ <https://github.com/wimvanderbauwheide/OpenCLIntegration>.

² <https://www.openacc.org/>.

³ <https://www.openmp.org/>.

⁴ <http://www.rosecompiler.org/index.html>.

⁵ <http://fortran-parser.sourceforge.net/>.

⁶ <https://hackage.haskell.org/package/language-fortran>.

⁷ <http://pips4u.org/>.

port FORTRAN 77 but does not support the refactorings we propose. Support for autotrialization via OpenCL was promised [7] but has not yet materialized. For completeness we mention the commercial solutions plusFort⁸ and VAST/77to90⁹ which both can refactor common blocks into modules but not into procedure arguments.

4. Our goal and approach

FORTRAN 77 code is often computationally efficient, and programmer efficient in terms of allowing the programmer to quickly write code and not be too strict about it. As a result it becomes very difficult to maintain and port. Our goal is that the refactored code should meet the following requirements:

4.1. Modern, maintainable and extensible

FORTRAN 77 was designed with very different requirements from today's languages, notably in terms of avoiding bugs. It is said that C gives you enough rope to hang yourself. If that is so then FORTRAN 77 provides the scaffold as well. Specific features that are unacceptable in a modern language are:

- Implicit typing, i.e. an undeclared variable gets a type based on its starting letter. This may be very convenient for the programmer but makes the program very hard to debug and maintain. Our compiler makes all types explicit (`implicit none`).
- No indication of the intended access of subroutine arguments: in FORTRAN 77 it is not possible to tell if an argument will be used read-only, write-only or read-write. This is again problematic for debugging and maintenance of code. Our compiler infers the intent for all subroutine and function arguments.
- In FORTRAN 77, procedures defined in a different source file are not identified as such. For extensibility as well as for maintainability, a module system is essential. Our compiler converts all non-program code units into modules which are used with an explicit `export (only)` declaration.

There are several more refactorings that our compiler applies, such as rewriting label-based loops as do-loops etc, but they are less important for this paper.

4.2. Accelerator-ready

As discussed in Section 2, the common feature of the vast majority of current accelerators is that they have a separate memory space, usually physically separate from the host memory. Furthermore, the common offload model is to create a "kernel" subroutine (either explicitly or implicitly) which is run on the accelerator device. Consequently, it is crucial to separate the memory spaces of the kernel and the host program.

- FORTRAN 77 programs make liberal use of global variables through "common" blocks. Our compiler converts these common block variables into subroutine arguments across the *complete* call tree of the program. Although refactoring of common blocks has been reported for some of the other projects, to our knowledge our compiler is the first to perform this refactoring across multiple nested procedure calls, potentially in different source code units.

⁸ <http://www.polyhedron.com/pf-plusfort0html>.

⁹ <http://www.crescentbaysoftware.com/compilertech.html>.

4.3. Automatic parallelization and acceleration

Our ultimate goal is to convert legacy FORTRAN 77 code into parallel code so that the computation can be accelerated using OpenCL. We use a three-step process:

First, the above refactorings¹⁰ result in a modern, maintainable, extensible and accelerator-ready Fortran 95 codebase. This is an excellent starting point for many of the other existing tools, for example the generated code can now easily be parallelized using OpenMP or OpenACC annotations, or further refactored if required using e.g. Photran or PIPS. However, we want to provide the user with an end-to-end solution that does not require any annotations.

The second step in our process is to identify data-level parallelism present in the code in the form of *maps* and *folds* (i.e. loops without dependencies and reductions). The terms *map* and *fold* are taken from functional programming and refer to ways of performing a given operation on all elements of a list. Broadly speaking these constructs are equivalent to loop nests with and without dependencies, and as Fortran is loop-based, our analysis is indeed an analysis of loops and dependencies. However, our internal representation uses the functional programming model where *map* and *fold* are functions operating on other functions (i.e. they are higher-order functions), the latter being extracted from the bodies of the loops. Thus we raise the abstraction level of our representation and make it independent of both the original code and the final code to be generated. We apply a number of rewrite rules for map- and fold-based functional programs (broadly speaking equivalent to loop fusion or fission) to optimise the code.

The third step is to generate OpenCL host and device code from the parallelized code. Because of the high abstraction level of our internal representation, we could easily generate OpenMP or OpenACC annotations, CUDA or Maxeler's MaxJ language used to program FPGAs. Our compiler¹¹ also minimizes the data transfer between the host and the accelerator by eliminating redundant transfers. This includes determining which transfers need to be made only once in the run of the program.

5. Code transformation validation

To assess the correctness and capability of our refactoring compiler, we used the NIST (US National Institute of Standards and Technology) FORTRAN 78 test suite,¹² which aims to validate adherence to the ANSI X3.9-1978 (FORTRAN 77) standard. We used a version with some minor changes:¹³ All files are properly formed; a non standard conforming FORMAT statement has been fixed in test file FM110.f; Hollerith strings in FORMAT statements have been converted to quoted strings. This test suite comprises about three thousand tests organized into 192 files. We skipped a number of tests because they test features that our compiler does not support. In particular, we skipped tests that use spaces in variable names and keywords (3 files, 23 tests) and tests for corner cases of common blocks and block data (2 files, 37 + 16 tests). After skipping these types of tests, 2867 tests remain, in total 187 files for which refactored code is generated. The test bench driver provided in the archive skips another 8 tests because they relate to features deleted in Fortran 95. In total the test suite contains 72,473 lines of code (excluding comments). Two test files contain tests that fail in gfortran 4.9 (3 tests in total).

Our compiler successfully generates refactored code for all tests, and the refactored code compiles correctly and passes all tests (2864 tests in total).

¹⁰ <https://github.com/wimvanderbauwheide/RefactorF4ACC>.

¹¹ <https://github.com/wimvanderbauwheide/AutoParallel-Fortran>.

¹² http://www.itl.nist.gov/div897/ctg/fortran_form.htm.

¹³ http://www.fortran-2000.com/ArnaudRecipes/fcv21_f95.html.

Furthermore, we tested the compiler on a simple 2-D shallow water model from [8] (188 loc) and on four real-word simulation models: the Large Eddy Simulator for Urban Flows,¹⁴ a high-resolution turbulent flow model [3] (1391 loc); the shallow water component of Gmodel¹⁵, an ocean model[9] (1533 loc); Flexpart-WRF,¹⁶ a version of the Flexpart particle dispersion simulator[10] that takes input data from WRF (13,829 loc); and the Linear Baroclinic Model,¹⁷ an atmospheric climate model [11] (39,336 loc).

Each of these models has a different coding style, specifically in terms of the use of common blocks, include files etc., that affect the refactoring process. All of these codes are refactored fully automatically without changes to the original code and build and run correctly. The performance of the original and refactored code is the same in all cases.

6. Automatic parallelization evaluation

In this section we show the performance of the automatically generated OpenCL code compared to the best achievable performance of the unmodified original code. We show that the automatically generated OpenCL code can perform as well as hand-ported OpenCL code.

6.1. Experimental setup

To evaluate the automatic parallelization and OpenCL code generation we used following experimental setup: the host platform is an Intel Xeon CPU E5-2620@2.00 GHz, a 6-core CPU with hyperthreading (12 threads), AVX, 32GB RAM, and 15MB cache; the GPU is an NVIDIA GeForce GTX TITAN, 980 MHz, 15 compute units, 16GB RAM. We used OpenCL 1.1 via the CUDA 6.5.14 SDK. The original UFLES code on CPU (reference) was compiled with gfortran 4.8.2 with following flags for auto-vectorization and auto-parallelization:-Ofast -floop-parallelize-all -ftree-parallelize-loops=12 -fopenmp -pthread. Auto-parallelization provides only 4% speed-up because the most time-consuming loops are not parallelized. Our compiler auto-parallelizes all loop nests in the code base and produces a complete OpenCL-enabled code base that runs on GPU and CPU.

6.2. Test case 1: 2-D Shallow Water model

As a first test case for the validation of our automatic parallelization approach we used the 2-D Shallow Water model from the textbook [8] by Kaempf. This very simple model consists of a time loop which calls two subroutines, a predictor (*dyn*) and a first-order Shapiro filter (*shapiro*), before updating the velocity. Our compiler automatically transforms this code into three map-style kernels.

The results shown in Fig. 2 are for domain size of 500×500 , 1000×1000 , and 2000×2000 for 10,000 time steps. This is a high-resolution simulation with spatial resolution of 1 m and a time step of 0.01 s. The automatically generated code running on GPU is up to 9x faster than the original code. This is the same performance as obtained by manual porting of the code to OpenCL.

6.3. Test case 2: Large Eddy Simulator for Urban Flows (UFLES)

As a more comprehensive test case we used the Large Eddy Simulator for Urban Flows (UFLES) developed by Prof. Takemi at the Disaster Prevention Research Institute of Kyoto University and

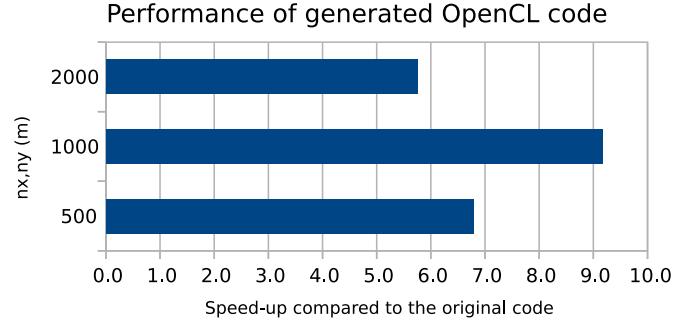


Fig. 2. Speed-up on NVIDIA GeForce GTX TITAN GPU compared to the original code.

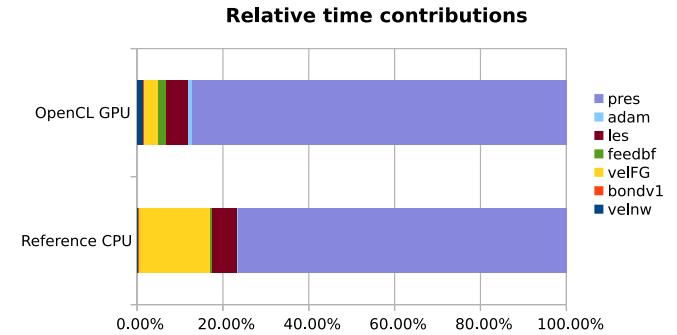


Fig. 3. Breakdown of time contribution per subroutine.

Dr. Nakayama of the Japan Atomic Energy Agency [12]. This simulator generates turbulent flows by using mesoscale meteorological simulations. It explicitly represents the urban surface geometry using GIS data and is used to conduct building-resolving large-eddy simulations of boundary-layer flows over urban areas under realistic meteorological conditions. The simulator essentially solves the Poisson equation for the pressure using Successive Over-Relaxation and integrates the force fields using the Adams–Bashforth algorithm.

6.3.1. Functional code structure of UFLES

The UFLES main loop sequentially executes 7 subroutines consecutively for each simulation time step:

- velnw: Update velocity for current time step
- bondv1: Calculate boundary conditions (initial wind profile, inflow, outflow)
- velfg: Calculate the body force
- feedbf: Calculation of building effects (Goldstein damping model)
- les: Calculation of viscosity terms (Smagorinsky model)
- adam: Adams-Bashforth time integration
- press: Solving of Poisson equation using SOR (iterative solver)

Our compiler automatically transforms this code into 29 map-style kernels and 4 reduction kernels.

6.3.2. OpenCL UFLES Results

All results shown in Figs. 2–4 and are for a domain size of $300 \times 300 \times 90$, with the number of SOR iterations set to 50. This is a realistic use case of the UFLES covering an area of 1.2 km \times 1.2 km. A simulation time step represents 0.025 s of actual time.

Fig. 3 shows the breakdown of relative run time contributions per subroutine. We can see that the *pres* subroutine which contains the SOR iterative loop dominates the run time. On the GPU, this routine accounts for almost 90% of the run time. Fig. 4 shows

¹⁴ <https://github.com/wimvanderbauwheide/LES>.

¹⁵ <http://www.sciamachy-validation.org/research/CKO/gmodel.html>.

¹⁶ https://github.com/sajinh/flx_wrf2.

¹⁷ <http://ccsr.aori.u-tokyo.ac.jp/~hiro/sub/lbm.html>.

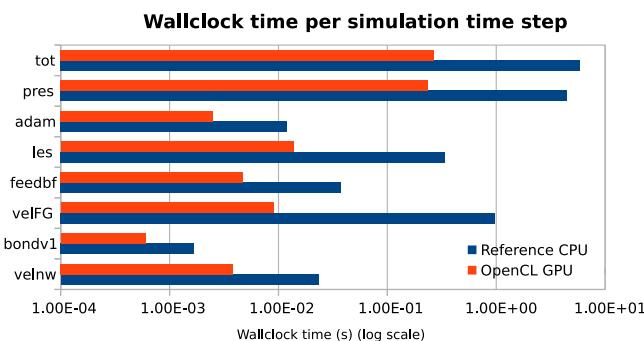


Fig. 4. Wall clock time.

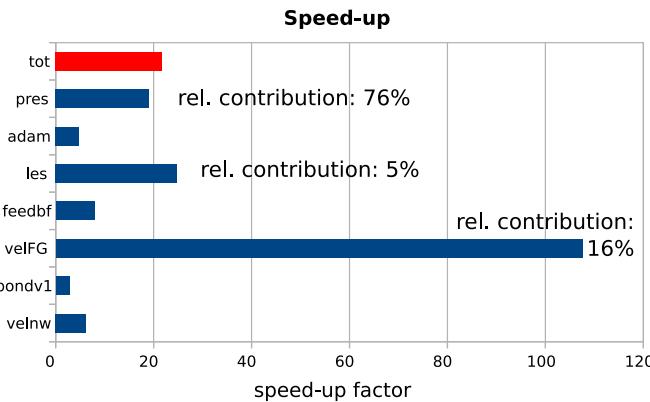


Fig. 5. Speed-up compared to the original code.

the total wall clock time and wall clock times for each subroutine on CPU and GPU. Note that the scale is logarithmic. The main observations are that the GPU code is faster for all subroutines but especially so for the *velFG* routine. Finally, Fig. 5 shows the total speed-up and the speed-up per subroutine. The speed-up of more than 100x for *velFG* is remarkable. This is because this routine performs a large amount of computations per point in the domain and each point is independent. Thus the GPU can optimally exploit the available parallelism. However, the total speed-up is entirely dominated by the iterative SOR solver, which is 20x faster on the GPU. Our auto-parallelized version achieves the same performance as the manually ported OpenCL version of the UFLES [3].

7. Discussion

The above results demonstrate that it is possible to automatically generate high-performance GPU code from FORTRAN 77 legacy code. All the compiler expects the programmer to do is annotate a region of the code for offloading. All subroutines in this region will be offloaded to the accelerator.

In practice there are some limitations. We have only presented two examples because the autoparallelizing compiler currently lacks a recursive inliner so that it only supports kernel subroutines that do not call other subroutines.

We use the term “domain specific” not in the sense of a particular branch of science but rather a class of models: in essence, we require the loop bounds to be static, i.e. known at compile time, in order to parallelize the loops. For the same reason, recursion is not supported; however, recursion is not supported by the ANSI X3.9-1978 (FORTRAN 77 standard). Furthermore, the current version of the compiler expects static array allocation, although this

is not a fundamental limitation and we are working on supporting dynamic allocation. The current OpenCL backend generates code that is optimized either for CPU or for GPU and we are actively working on generating optimized code for FPGAs.

8. Conclusion

We have developed a proof-of-concept compiler for OpenCL acceleration and auto-parallelization of domain-specific legacy FORTRAN 77 scientific code using whole-program analysis and source-to-source compilation. We have validated the code transformation performance of the compiler on the NIST FORTRAN78 test suite and a number of real-world codes; the automatic parallelization component has been tested on a 2-D Shallow Water model and on the Large Eddy Simulator for Urban Flows and produces a complete OpenCL-enabled code base that is 20x faster on GPU than the original code on CPU. Future work will focus on improving the compiler to extract more parallelism from the original code and improve the performance; and development of a complete FPGA back-end.

Acknowledgements

The authors acknowledge the support of the EPSRC (EP/L00058X/1).

Supplementary material

Supplementary material associated with this article can be found, in the online version, at doi:10.1016/j.compfluid.2018.06.005.

References

- [1] Tinetti FG, Méndez M. Fortran legacy software: source code update and possible parallelisation issues. In: ACM SIGPLAN fortran forum, 31. ACM; 2012. p. 5–22.
- [2] Stone JE, Gohara D, Shi G. OpenCL: a parallel programming standard for heterogeneous computing systems. Comput Sci Eng 2010;12(3):66–73.
- [3] Vanderbauwheide W, Takemi T. Twinned buffering: a simple and highly effective scheme for parallelization of successive over-relaxation on gpus and other accelerators. In: High performance computing & simulation (HPCS), 2015 int'l conf on. IEEE; 2015. p. 436–43.
- [4] Liao C, Quinlan DJ, Panas T, De Supinski BR. A rose-based openmp 3.0 research compiler supporting multiple runtime libraries. In: International workshop on OpenMP. Springer; 2010. p. 15–28.
- [5] Orchard D, Rice A. Upgrading fortran source code using automatic refactoring. In: Proceedings of the 2013 ACM workshop on workshop on refactoring tools, WRT'13. New York, NY, USA: ACM; 2013. p. 29–32. ISBN 978-1-4503-2604-9. doi:10.1145/2541348.2541356.
- [6] Overbeek J, Xanthos S, Johnson R, Foote B. Refactorings for fortran and high-performance computing. In: Proceedings of the second international workshop on Software engineering for high performance computing system applications. ACM; 2005. p. 37–9.
- [7] Amini M, Arcourt C, Coelho F, Creusillet B, Guelton S, Irigoin F, et al. Pips is not (just) polyhedral software adding gpu code generation in pips. In: First international workshop on polyhedral compilation techniques (IMPACT 2011) in conjunction with CGO 2011; 2011. p. 1–6.
- [8] Kämpf J. Ocean modelling for beginners: using open-source software. Springer Science & Business Media; 2009.
- [9] Burgers G, Balmaseda MA, Vossepoel FC, van Oldenborgh GJ, Van Leeuwen PJ. Balanced ocean-data assimilation near the equator. J Phys Oceanogr 2002;32(9):2509–19.
- [10] Brioude J, Arnold D, Stohl A, Cassiani M, Morton D, Seibert P, et al. The lagrangian particle dispersion model flexpart-wrf version 3.1. Geosci Model Dev 2013;6(6):1889–904.
- [11] Watanabe M, Jin F-F. A moist linear baroclinic model: coupled dynamical–convective response to el niño. J Clim 2003;16(8):1121–39.
- [12] Nakayama H, Takemi T, Nagai H. Les analysis of the aerodynamic surface properties for turbulent flows over building arrays with various geometries. J Appl Meteorol Climatol 2011;50(8):1692–712.