


S9624:

Performance Analysis of GPU-Accelerated Applications using the Roofline Model

GTC 2019, San Jose



Charlene Yang
Application Performance Specialist
NERSC, LBNL
cjyang@lbl.gov

Samuel Williams
Senior Staff Scientist
CRD, LBNL
swwilliams@lbl.gov




You just bought a \$10,000
throughput-optimized GPU!

Are you making good use of
your investment?

You could just run benchmarks

- Imagine a mix of benchmarks or kernels...
- GFLOP/s alone may not be particularly insightful
- Moreover, speedup relative to a Xeon may seem random



Making good use of your GPU?




1. Are you operating it in the throughput-limited regime?
 - Not sensitive to Amdahl effects
 - Not sensitive to D2H/H2D transfers
 - Not sensitive to launch overheads
 - Not sensitive to latencies

2. If in the throughput-limited regime, are you making good use of the GPU's **compute** and **bandwidth** capabilities?

The Roofline Model



- **Roofline Model** is a throughput-oriented performance model
- Premised on the interplay between FLOP/s, bandwidth, and reuse
- Tracks rates not times
- Independent of ISA and architecture (applies to CPUs, GPUs, Google TPUs, etc...)




<https://crd.lbl.gov/departments/computer-science/PAR/research/roofline>

Jouppi et al, “In-Datacenter Performance Analysis of a Tensor Processing Unit”, ISCA, 2017.

(DRAM) Roofline



- One could hope to always attain peak performance (GFLOP/s)
- However, finite locality (reuse) and bandwidth limit performance.
- Assume:
 - Idealized processor/caches
 - Cold start (data in DRAM)




$$\text{Time} = \max \left\{ \begin{array}{l} \text{\#FLOPs / Peak GFLOP/s} \\ \text{\#Bytes / Peak GB/s} \end{array} \right\}$$

(DRAM) Roofline



- One could hope to always attain peak performance (GFLOP/s)
- However, finite locality (reuse) and bandwidth limit performance.
- Assume:
 - Idealized processor/caches
 - Cold start (data in DRAM)



$$\text{GFLOP/s} = \min \left\{ \begin{array}{l} \text{Peak GFLOP/s} \\ \text{AI * Peak GB/s} \end{array} \right\}$$

Note, Arithmetic Intensity (AI) = FLOPs / Bytes (as presented to DRAM)

Arithmetic Intensity



- **Arithmetic Intensity** is the most important concept in Roofline.

- Measure of data locality (data reuse)
- Ratio of **Total FLOPs** performed to **Total Bytes** moved
- For the DRAM Roofline...
 - Total Bytes to/from DRAM and includes all cache and prefetcher effects
 - Can be **very different from total loads/stores** (bytes requested) due to cache reuse



U.S. DEPARTMENT OF
ENERGY


Office of
Science



(DRAM) Roofline




- Plot Roofline bound using Arithmetic Intensity as the x-axis
- **Log-log scale** makes it easy to doodle, extrapolate performance along Moore's Law, etc...
- Kernels with AI less than machine balance are ultimately DRAM bound (we'll refine this later...)



Example




- Consider 3 kernels (A,B,C)
 - calculate or measure the **Arithmetic Intensity** for each
 - Determine the Roofline intercept for each kernel
 - **kernels A and B are bound by memory bandwidth**
 - **kernel C is bound by peak FLOP/s**



Scaling to Future GPUs




- Imagine you run on a future GPU with twice the peak FLOPs...
 - **kernel C's performance could double**
 - ✖ **kernels A and B will be no faster**



Scaling to Future GPUs




- What if that future GPU also doubled its memory bandwidth...
 - **kernel A and B's performance could also double**




Why is Roofline Useful?

- Think back to our mix of loop nests where GFLOP/s alone wasn't useful...




Why is Roofline Useful?

- We can sort kernels by AI ...




Why is Roofline Useful?

- We can sort kernels by AI ...
- ... and compare performance relative to machine capabilities



Why is Roofline Useful?


- Kernels near the roofline are making good use of computational resources...
 - kernels can have low performance (GFLOP/s), but make good use of a machine
 - kernels can have high performance (GFLOP/s), but make poor use of a machine



Can Performance Be Below Roofline?




- Analogous to asking whether one can always attain either...
 - Peak Bandwidth
 - Peak GFLOP/s
- **Sure, there can be other performance bottlenecks...**
 - Cache bandwidth / locality
 - Lack of FMA / tensor instructions
 - Thread divergence / predication
 - Too many non-FP instructions
 - ...



Cache Effects...




- Hierarchical Roofline Model
- Construct superposition of Rooflines...
 - Measure AI and bandwidth for each level of memory/cache
 - Loop nests will have multiple AI's and multiple performance bounds...
 - **... but performance is ultimately the minimum of these bounds.**



Cache Effects...




- Hierarchical Roofline Model
- Construct superposition of Rooflines...
 - Measure AI and bandwidth for each level of memory/cache
 - Loop nests will have multiple AI's and multiple performance bounds...
 - ... but performance is ultimately the minimum of these bounds.
- Extend to other memories...
 - L1 / Shared
 - System



Insights – Exploiting Caches




- Widely separated Arithmetic Intensities indicate high reuse in the cache



Insights – Exploiting Caches



- Widely separated Arithmetic Intensities indicate high reuse in the cache
- Similar Arithmetic Intensities indicate effectively no cache reuse (**== streaming**)
- As one changes problem size, L2 and DRAM arithmetic intensities can behave very differently



Failure to Exploit CISC Instructions




- Death of Moore's Law is motivating a return of Complex Instruction Set Computing (CISC)
 - Modern CPUs and GPUs are increasingly reliant on special (fused) instructions that perform multiple operations.
 - FMA (Fused Multiply Add): $z=a*x+y$... z,x,y are vectors or scalars
 - 4FMA (quad FMA): $z=A*x+z$... A is a FP32 matrix; x,z are vectors
 - HMMA (Tensor Core): $Z=AB+C$... Z,A,B,C are FP16 matrices
 - ...
- **Performance is now a weighted average of Mul/Add, FMA, and HMMA operations.**

Failure to Exploit CISC Instructions




- Total lack of FMA reduces Volta performance by 2x...
 - **creates ADD.f64 ceiling**
- In reality, applications are a mix of FMA.f64, ADD.f64, and MUL.f64...
 - Performance is a weighted average
 - **Produces a partial FMA ceiling that bounds kernel performance**




Failure to Exploit CISC Instructions



- On Volta, Tensor cores provide **125 TFLOPs** of FP16 performance (vs. 15 for FP32)
- However, kernels/apps will mix HMMA with FMA, MULs, ADDs, ...
 - **A few non-HMMA operations can quickly limit Tensor core performance**






Using Roofline To Drive Optimization

Driving Performance Optimization




- Broadly speaking, there are three approaches to improving performance:



Driving Performance Optimization




- Broadly speaking, there are three approaches to improving performance:
- **Maximize SM performance (e.g. minimize predication)**



Driving Performance Optimization




- Broadly speaking, there are three approaches to improving performance:
- Maximize SM performance (e.g. minimize predication)
- **Maximize memory bandwidth (e.g. avoid pathological memory access patterns)**




Driving Performance Optimization



- Broadly speaking, there are three approaches to improving performance:
- Maximize SM performance (e.g. minimize predication)
- Maximize memory bandwidth (e.g. avoid pathological memory access patterns)
- **Minimize data movement
(i.e. exploit reuse)**





Estimating Arithmetic Intensity


DRAM vs L1 Arithmetic Intensity



- Consider a 7-point constant coefficient stencil...

- 7 FLOPs
- 8 memory references (7 reads, 1 store) per point
- **AI = 0.11 FLOPs per byte (L1)**

```
#pragma omp parallel for
for(k=1;k<dim+1;k++){
    for(j=1;j<dim+1;j++){
        for(i=1;i<dim+1;i++){
            new[k][j][i] = -6.0*old[k][j][i]
                + old[k][j][i-1]
                + old[k][j][i+1]
                + old[k][j-1][i]
                + old[k][j+1][i]
                + old[k-1][j][i]
                + old[k+1][j][i];
        }
    }
}
```




DRAM vs L1 Arithmetic Intensity



- Consider a 7-point constant coefficient stencil...
 - 7 FLOPs
 - 8 memory references (7 reads, 1 store) per point
 - Cache can filter all but 1 read and 1 write per point
 - **AI = 0.44 FLOPs per byte**

```
#pragma omp parallel for
for(k=1;k<dim+1;k++){
    for(j=1;j<dim+1;j++){
        for(i=1,i<dim+1;i++){
            new[k][j][i] = -6.0*old[k][j][i]
                + old[k][j][i-1]
                + old[k][j][i+1]
                + old[k][j-1][i]
                + old[k][j+1][i]
                + old[k-1][j][i]
                + old[k+1][j][i]
        }
    }
}
```





DRAM vs L1 Arithmetic Intensity



- Consider a 7-point constant coefficient stencil...
 - 7 FLOPs
 - 8 memory references (7 reads, 1 store) per point
 - Cache can filter all but 1 read and 1 write per point
 - **AI = 0.44 FLOPs per byte == memory bound**

```
#pragma omp parallel for
for(k=1;k<dim+1;k++){
  for(j=1;j<dim+1;j++){
    for(i=1;i<dim+1;i++){
      new[k][j][i] = -6.0*old[k][j][i]
                    + old[k][j][i-1]
                    + old[k][j][i+1]
                    + old[k][j-1][i]
                    + old[k][j+1][i]
                    + old[k-1][j][i]
                    + old[k+1][j][i];
    }
  }
}
```





Collecting Roofline Data with nvprof

General Roofline Data Collection



Most kernels are more complicated than the 7-point stencil...

General Roofline Data Collection



Most kernels are more complicated than the 7-point stencil...

How do we measure the total number of FLOPs?

How do we measure the total number of bytes moved (read/write, L1/L2/HBM)?

How do we measure the runtime for each kernel?

How do we know the peak bandwidth (L1/L2/HBM) and the peak FLOP/s for the architecture?

General Roofline Data Collection




Most kernels are more complicated than the 7-point stencil...

How do we measure the total number of FLOPs?

How do we measure the total number of bytes moved (read/write, L1/L2/HBM)?

How do we measure the runtime for each kernel?




How do we know the peak bandwidth (L1/L2/HBM) and the peak FLOP/s for the architecture?



Step 1. Collect Roofline Ceilings



- **Empirical Roofline Toolkit (ERT)**
 - Different than the architecture specs, **MORE REALISTIC**
 - Reflects **actual** execution environment (power constraints, etc)
 - Sweeps through a range of configurations, and **statistically stable**
 - Data elements per thread
 - FLOPs per data element
 - Threadblocks/threads
 - Trials per dataset
 - *etc*



ERT Configuration



Kernel.c

- actual compute
- customizable

Driver.c

- setup
- call kernels
- loop over parameters

config script

- set up ranges of parameters

job script

- submit the job and run it

ERT Output




roofline.json

```
"gbytes": {  
    "data": [  
        [  
            "L1",  
            2996.82  
        ],  
        [  
            "DRAM",  
            828.83  
        ]  
    ],  
},
```

```
"gflops": {  
    "data": [  
        [  
            "GFLOPs",  
            7068.90  
        ]  
    ],  
},
```

roofline.ps



ERT Output




roofline.json

```
"gbytes": {  
    "data": [  
        [  
            "L1",  
            2996.82  
        ],  
        [  
            "DRAM",  
            828.83  
        ]  
    ],  
}
```

```
"gflops": {  
    "data": [  
        [  
            "GFLOPs",  
            7068.90  
        ]  
    ],  
}
```

roofline.ps



ERT Output




roofline.json

```
"gbytes": {  
    "data": [  
        [  
            "L1",  
            2996.82  
        ],  
        [  
            "DRAM",  
            828.83  
        ]  
    ],  
}
```

```
"gflops": {  
    "data": [  
        [  
            "GFLOPs",  
            7068.90  
        ]  
    ],  
}
```


roofline.ps




Discrepancy Empirical vs. Theoretical




- Theoretical FP64 **compute** ceilings on V100:
 - FMA: $80 \text{ SMs} \times 32 \text{ FP64 cores} \times 1.53 \text{ GHz} \times 2 = 7.83 \text{ TFLOP/s}$
 - no FMA: $80 \text{ SMs} \times 32 \text{ FP64 cores} \times 1.53 \text{ GHz} = 3.92 \text{ TFLOP/s}$
- Theoretical **memory** bandwidths on V100:
 - HBM: 900 GB/s
 - L2: ~4.1 TB/s
 - L1: ~14 TB/s
- You may never achieve 7.8 TFLOP/s**
- You may be closer to the ceiling than you think you are**




Step 2. Collect Application Performance



Step 2. Collect Application Performance



Step 2. Collect Application Performance



Require three raw measurements:

- **Runtime**
- **FLOPs**
- **Bytes (on each cache level)**

to calculate AI and GFLOP/s:

$$\text{Arithmetic Intensity} = \frac{\text{nvprof FLOPs}}{\text{nvprof Data Movement}}$$

(FLOPs/Byte)

$$\text{Performance} = \frac{\text{nvprof FLOPs}}{\text{Runtime}}$$

(GFLOP/s)

Collect Application Performance



- Runtime:
 - Time per invocation of a kernel

```
nvprof --print-gpu-trace ./application
```
 - Average time over multiple invocations

```
nvprof --print-gpu-summary ./application
```
 - Same kernel with different input parameters are grouped separately
- FLOPs:
 - Predication aware and complex-operation aware (such as divides)
 - ```
nvprof --kernels 'kernel_name' --metrics 'flop_count_xx' ./application
```
  - e.g. `flop_count_{dp/dp_add/dp_mul/dp_fma, sp*, hp*}`

# Collect Application Performance



- Bytes for different cache levels in order to construct hierarchical Roofline:
  - Bytes = (read transactions + write transactions) x transaction size
  - `nvprof --kernels 'kernel_name' --metrics 'metric_name'`  
`./application`

| Level              | Metrics                                                                                                                                                                                                      | Transaction Size |
|--------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
| First Level Cache* | <code>gld_transactions, gst_transactions, atomic_transactions,</code><br><code>local_load_transactions, local_store_transactions,</code><br><code>shared_load_transactions, shared_store_transactions</code> | 32B              |
| Second Level Cache | <code>l2_read_transactions, l2_write_transactions</code>                                                                                                                                                     | 32B              |
| Device Memory      | <code>dram_read_transactions, dram_write_transactions</code>                                                                                                                                                 | 32B              |
| System Memory      | <code>system_read_transactions, system_write_transactions</code>                                                                                                                                             | 32B              |

- Note: surface and texture transactions are ignored here for simplicity (HPC applications)

# Example Output

```
[cjyang@voltar source]$ nvprof --kernels "1:7:smooth_kernel:1" --metrics flop_count_dp --metrics gld_transactions --metrics gst_transactions --metrics l2_read_transactions --metrics l2_write_transactions --metrics dram_read_transactions --metrics dram_write_transactions --metrics sysmem_read_bytes --metrics sysmem_write_bytes ./hpgmg-fv-fp 5 8
```

- Export to CSV: **--csv -o nvprof.out**

context : stream : kernel : invocation

| Invocations                                                                                                          | Metric Name             | Metric Description                          | Min      | Max      | Avg      |
|----------------------------------------------------------------------------------------------------------------------|-------------------------|---------------------------------------------|----------|----------|----------|
| Device "Tesla V100-PCIE-16GB (0)"                                                                                    |                         |                                             |          |          |          |
| Kernel: void smooth_kernel<int=6, int=32, int=4, int=8>(level_type, int, int, double, double, int, double*, double*) | flop_count_dp           | Floating Point Operations(Double Precision) | 30277632 | 30277632 | 30277632 |
|                                                                                                                      | gld_transactions        | Global Load Transactions                    | 4280320  | 4280320  | 4280320  |
|                                                                                                                      | gst_transactions        | Global Store Transactions                   | 73728    | 73728    | 73728    |
|                                                                                                                      | l2_read_transactions    | L2 Read Transactions                        | 890596   | 890596   | 890596   |
|                                                                                                                      | l2_write_transactions   | L2 Write Transactions                       | 85927    | 85927    | 85927    |
|                                                                                                                      | dram_read_transactions  | Device Memory Read Transactions             | 702911   | 702911   | 702911   |
|                                                                                                                      | dram_write_transactions | Device Memory Write Transactions            | 151487   | 151487   | 151487   |
|                                                                                                                      | sysmem_read_bytes       | System Memory Read Bytes                    | 0        | 0        | 0        |
|                                                                                                                      | sysmem_write_bytes      | System Memory Write Bytes                   | 160      | 160      | 160      |

# Step 3. Plot Roofline with Python



- Calculate Arithmetic Intensity and GFLOP/s performance
  - x coordinate: Arithmetic Intensity
  - y coordinate: GFLOP/s performance

$$\text{Performance} = \frac{\textit{nvprof FLOPs}}{\text{Runtime}} \quad (\text{GFLOP/s}), \quad \text{Arithmetic Intensity} = \frac{\textit{nvprof FLOPs}}{\textit{nvprof Data Movement}} \quad (\text{FLOPs/Byte})$$

- Plot Roofline with Python Matplotlib
  - Example scripts:
  - <https://github.com/cyanguwa/nersc-roofline/tree/master/Plotting>
  - Tweak as needed for more complex Rooflines

# Plot Roofline with Python




- Quick example: `plot_roofline.py data.txt`
- Accepts space-delimited list for values
- Use quotes to separate names/labels

`data.txt`

```
all data is space delimited
memroofs 14336.0 2996.8 828.758
mem_roof_names 'L1' 'L2' 'HBM'
comproofs 7068.86 3535.79
comp_roof_names 'FMA' 'No-FMA'

omit the following if only plotting roofs
AI: arithmetic intensity; GFLOPs: performance
AI 0.87 2.25 2.58
GFLOPs 2085.756683
labels 'Kernel'
```



# Recap: Methodology to Construct Roofline



## 1. Collect Roofline ceilings

- ERT: <https://bitbucket.org/berkeleylab/cs-roofline-toolkit>
- compute (FMA/no FMA) and bandwidth (DRAM, L2, ...)

## 2. Collect application performance

- nvprof: `--metrics`, `--events`, `--print-gpu-trace`
- FLOPs, bytes (DRAM, L2, ...), runtime

## 3. Plot Roofline with Python Matplotlib

- arithmetic intensity, GFLOP/s performance, ceilings
- example scripts: <https://github.com/cyanguwa/nersc-roofline>

# Recap: Methodology to Construct Roofline



## 1. Collect Roofline ceilings

- ERT: <https://bitbucket.org/berkeleylab/cs-roofline-toolkit>
- **compute** (FMA/no FMA) and **bandwidth** (DRAM, L2, ...)

## 2. Collect application performance

- nvprof: **--metrics**, **--events**, **--print-gpu-trace**
- FLOPs, bytes (DRAM, L2, ...), runtime

## 3. Plot Roofline with Python Matplotlib

- arithmetic intensity, GFLOP/s performance, ceilings
- example scripts: <https://github.com/cyanguwa/nersc-roofline>

# Recap: Methodology to Construct Roofline



## 1. Collect Roofline ceilings

- ERT: <https://bitbucket.org/berkeleylab/cs-roofline-toolkit>
- **compute** (FMA/no FMA) and **bandwidth** (DRAM, L2, ...)

## 2. Collect application performance

- nvprof: **--metrics**, **--events**, **--print-gpu-trace**
- **FLOPs**, **bytes** (DRAM, L2, ...), **runtime**

## 3. Plot Roofline with Python Matplotlib

- arithmetic intensity, GFLOP/s performance, ceilings
- example scripts: <https://github.com/cyanguwa/nersc-roofline>

# Recap: Methodology to Construct Roofline



## 1. Collect Roofline ceilings


- ERT: <https://bitbucket.org/berkeleylab/cs-roofline-toolkit>
- **compute** (FMA/no FMA) and **bandwidth** (DRAM, L2, ...)

## 2. Collect application performance

- nvprof: **--metrics**, **--events**, **--print-gpu-trace**
- **FLOPs**, **bytes** (DRAM, L2, ...), **runtime**

## 3. Plot Roofline with Python Matplotlib

- **arithmetic intensity**, **GFLOP/s** performance, **ceilings**
- example scripts: <https://github.com/cyanguwa/nersc-roofline>



# Roofline Analysis with Use Cases

# Code Example 1: GPP



- GPP (General Plasmon Pole) kernel from BerkeleyGW (Material Science)
- <https://github.com/cyanguwa/BerkeleyGW-GPP>
- Medium problem size: 512 2 32768 20
- Tensor-contraction, abundant parallelism, large reductions
- Low FMA counts, divides, complex double data type, HBM data 1.5GB

## Pseudo Code

```
do band = 1, nbands #blockIdx.x
 do igr = 1, ngpown #blockIdx.y
 do ig = 1, ncouls #threadIdx.x
 do iw = 1, nw #unrolled
 compute; reductions
```

# Code Example 1: GPP



- Three experiments:

|                                  |                                                                       |
|----------------------------------|-----------------------------------------------------------------------|
| Vary <code>nw</code> from 1 to 6 | To study impact of <b>varying Arithmetic Intensity</b> on performance |
| Compile w/wo FMA                 | To study impact of <b>instruction mix</b> on performance              |
| Stride <code>ig</code> loop      | To study impact of <b>suboptimal memory coalescing</b> on performance |

- Note that `nvprof` has already taken care of
  - Appropriate counting of FLOPs for complex instructions
    - div, exp, log and sin/cos should be counted as multiple FLOPs rather than 1
  - Appropriate counting of FLOPs for predicated-out threads
    - FLOPs are only counted on non-predicated threads

# Code Example 1: GPP

- Highly parameterizable
  1. Varying **nw** from 1 to 6 to increase arithmetic intensity
    - FLOPs increases, but data movement stays (at least for HBM)

## Pseudo Code

```
do band = 1, nbands #blockIdx.x
 do igr = 1, ngpown #blockIdx.y
 do ig = 1, ncouls #threadsIdx.x
 do iw = 1, nw #unrolled
 compute; reductions
```

2. Compiling with and without FMA
  - **-fmad=true/false**

# Code Example 1: GPP


- Highly parameterizable
  - 3. Striding **ig** loop to analyze impact of suboptimal memory coalescing
    - Split **ig** loop to two loops and place the ‘blocking’ loop outside

## Pseudo Code

```

do band = 1, nbands
 do igr = 1, ngpown
 do igs = 0, stride - 1
 do ig = 1, ncouls/stride #threadIdx.x
 do iw = 1, nw #unrolled
 compute; reductions


```



# Code Example 1: GPP




- **Experiments 1:** study the impact of varying AI on performance
- HBM Roofline, i.e. bytes are HBM bytes
  - AI increases as `nw` grows
  - GPP moves from a bandwidth bound region to a compute bound region
- Roofline captures the change in AI



# Code Example 1: GPP




- **Experiments 1 & 2:** study the impact of instruction mix on performance
- HBM Roofline, i.e. bytes are HBM bytes
  - No-FMA performance converges to the no-FMA ceiling, but FMA performance is still far from the FMA ceiling
  - Not reaching FMA ceiling due to lack of FMA instructions
- Roofline captures effects of instruction mix



# Code Example 1: GPP


- **Experiments 1 & 2:** study the impact of instruction mix on performance
- At  $nw=6$ , GPP has  $\alpha = \frac{\text{FMA FP64 instr.}}{\text{FMA FP64 instr.} + \text{non-FMA FP64 instr.}} = 60\%$  of FMA instructions
- Expected performance is  $\beta = \frac{\alpha \times 2 + (1 - \alpha)}{2} = 80\%$  of compute peak.  
But at  $nw=6$ , GPP is only achieving **66%**.
- Other FP/non-FP instructions may be taking up the instruction issue/execution pipeline
- Partial Roofline can show you the headroom



# Code Example 1: GPP




- **Experiments 1 & 2:** What else is going on?
- Hierarchical Roofline, i.e. bytes are HBM, L2 and unified L1 cache bytes
  - GPP is HBM bound at low **nw**'s and compute bound at high **nw**'s
  - FLOPs  $\propto$  **nw**
  - HBM bytes: constant
  - L2 bytes: increasing at  $\alpha > 1$
  - L1 bytes: constant
  - Spike in L2 curve at **nw=2, 3**
- Hierarchical Roofline captures more details about cache locality



# Code Example 1: GPP



- **Experiment 3:** study the effects of suboptimal memory coalescing
  - **nw=6**
- Hierarchical Roofline, i.e. bytes are HBM, L2 and unified L1 cache bytes
  - L1/L2 bytes doubles from stride 1 to 2, but stays almost constant afterwards
  - at **nw=6**, GPP moves from compute bound to bandwidth bound
  - Eventually all dots converge to HBM
- Roofline captures effects of memory coalescing



# Code Example 2: HPGMG





- HPGMG (High-performance Geometric Multigrid) from Adaptive Mesh Refinement codes
- <https://bitbucket.org/nsakharnykh/hpgmg-cuda>
- Stencil code, F-cycles and V-cycles, GSRB smoother kernel (Gauss-Seidel Red-Black)



# Code Example 2: HPGMG

- Hybrid GPU and CPU code
  - Example: `hpgmg-fv 7 8`
  - $128^3$  box x 8, Level 5-8 run on GPU, Level 1-4 on CPU
- Three versions of GSRB kernel
  - GSRB\_FP, GSRB\_BRANCH, GSRB\_STRIDE2



# Code Example 2: HPGMG




**GSRB\_FPs**

```
for(int k=klo; k<(klo+kdim) ; k++) {
 const int ijk = i + j*jStride + k*kStride;
 const double * __restrict__ RedBlack =
 level.RedBlack_FP + ghosts*(1+jStride)
 + ((k^color000)&1)*kStride;
 const double Ax = apply_op_ijk();
 const double lambda = Dinv_ijk();
 const int ij = i + j*jStride;
 xo[ijk] = X(ijk) + RedBlack[ij]*lambda*(rhs[ijk]-Ax);
}
```



8 elements

Sweep




8 threads

# Code Example 2: HPGMG



## GSRB\_FP

- Hierarchical Roofline, i.e. bytes are HBM, L2 and unified L1 cache bytes
- Highly bandwidth bound, inherent to stencil codes
- From Level 5 to Level 8:
  - AI slightly increases due to better Surface: Volume ratio
  - More HBM bound as more data is read in
- Roofline captures computational characteristics of the algorithm



# Code Example 2: HPGMG




## GSRB\_FP

```
for(int k=klo; k<(klo+kdim); k++) {
 const int ijk = i + j*jStride + k*kStride;
 const double * __restrict__ RedBlack =
 level.RedBlack_FP + ghosts*(1+jStride)
 +(k^color000)&1)*kStride;
 const double Ax = apply_op_ijk();
 const double lambda = Dinv_ijk();
 const int ij = i + j*jStride;
 xo[ijk] = X(ijk) + RedBlack[ij]*lambda*(rhs[ijk]-Ax);
}
```



## GSRB\_BRANCH

```
for(int k=klo; k<klo+kdim; k++) {
 const int ijk = i + j*jStride + k*kStride;
 if(((i^j^k^color000^1)&1))
 const double Ax = apply_op_ijk();
 const double lambda = Dinv_ijk();
 xo[ijk] = X(ijk) + lambda*(rhs[ijk]-Ax);
 else{
 xo[ijk] = X(ijk);
 }
}
```



8 elements



8 elements

Sweep



8 threads



8 threads



- GSRB\_BRANCH has half the FLOPs as GSRB\_FP but the same HBM/L1/L2 bytes

# Code Example 2: HPGMG



## GSRB\_FP vs. GSRB\_BRANCH

- FLOPs halves, bytes doesn't change, thus AI halves and GFLOP/s halves
- Runtime is comparable even though GFLOP/s has halved
- Same number of threads occupied, only with half predicated in GSRB\_BRANCH




# Code Example 2: HPGMG



```
GSRB_STRIDE2

for(int k=klo; k<klo+kdim; k++) {
 i = ilo +!((ilo^j^k^color000)&1) + threadIdx.x*2;
 if(i < ilo+idim) {
 const int ijk = i + j*jStride + k*kStride;
 xo[ijk] = X(ijk);
 }
 i = ilo + ((ilo^j^k^color000)&1) + threadIdx.x*2;
 if(i < ilo+idim) {
 const int ijk = i + j*jStride + k*kStride;
 const double Ax = apply_op_ijk();
 const double lambda = Dinv_ijk();
 xo[ijk] = X(ijk) + lambda*(rhs[ijk]-Ax);
 }
}
```





- GSRB\_STRIDE2 should have the same FLOPs as GSRB\_BRANCH, but same bytes?  
More writes than GSRB\_BRANCH?

# Code Example 2: HPGMG



## GSRB\_BRANCH vs. GSRB\_STRIDE2


- Extra writes in GSRB\_STRIDE2 cause more capacity misses in L2, leading to AI drop on L2 and DRAM, starting from Level 7 (data size  $\approx$  L2 cache size)
- Runtime almost doubled and GFLOP/s halved



# Conclusions



- Roofline can gracefully capture various aspects of application performance and architecture characteristics such as arithmetic intensity, instruction mix, memory coalescing and thread predication.
- The proposed methodology is effective in collecting machine characteristics and application data on NVIDIA GPUs to construct **hierarchical Roofline**.
- The Roofline model provides **insights** that profile:
  - identify the most immediate bottleneck
  - prioritize optimization efforts
  - tell you when you can stop



# Reference



- S. Williams, A. Waterman and D. Patterson, “Roofline: An insightful visual performance model for multicore architectures,” *Communications of the ACM*, vol. 52, no. 4, pp. 65–76, 2009
- Empirical Roofline Toolkit (ERT): <https://bitbucket.org/berkeleylab/cs-roofline-toolkit>
- Example scripts for plotting Roofline: <https://github.com/cyanguwa/nersc-roofline>
- General Plasmon Pole kernel: <https://github.com/cyanguwa/BerkeleyGW-GPP>
- HPGMG-CUDA kernel: <https://bitbucket.org/nsakharnykh/hpgmg-cuda>



# Acknowledgement

---

- This material is based upon work supported by the Advanced Scientific Computing Research Program in the U.S. Department of Energy, Office of Science, under Award Number DE-AC02-05CH11231.
- This material is based upon work supported by the DOE RAPIDS SciDAC Institute.
- This research used resources of the National Energy Research Scientific Computing Center (NERSC), which is supported by the Office of Science of the U.S. Department of Energy under contract DE-AC02- 05CH11231.



U.S. DEPARTMENT OF  
**ENERGY**

Office of  
Science



# Thank You!

